

winters_2025_brain_tumor_cnn_clasifcation_v0.2

March 7, 2025

- Wiley Winters
- MSDS 686 Deep Learning
- Week 7-8 Kaggle Project — Brain Tumor Classification
- 2025-MAR-09

0.1 Requirements

0.1.1 Required for 80%

Complete project on *kaggle.com* using the skills learned in the Deep Learning class. The following are required: - Show/plot sample images or data with labels - Include at least one of the following - Convolution - Max Pooling - Batch Normalization - Dropout - LSTM - TF-IDf - Use validation data - Evaluate model on test data

0.2 Additional for another 20%

- Use data augmentation
 - Use at least one of the following:
 - Kernels
 - Activation functions
 - Loss functions
 - Libraries
 - Methods
 - Learning rate optimization
 - Functional API model
 - Transfer learning with or without trainable parameters
 - Confusion matrix and / or ROC plots
 - Plots of accuracy/loss vs epochs
 - Show/plot sample incorrect prediction with labels and correct label
-

1.0 | Load Libraries and Packages

```
[1]: # General Imports
import numpy as np
```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os, logging, random
from datetime import datetime

# Data prep and model scoring
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score

# TensorFlow likes to display a lot of debug information
# on my home system
# I will squash the messages
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
logging.getLogger('tensorflow').setLevel(logging.FATAL)

# tensorflow and keras' API
import tensorflow as tf
from tensorflow import keras

# Model building
from tensorflow.keras import backend, optimizers, regularizers, models
from tensorflow.keras.layers import Dense, Dropout, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D, BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras.regularizers import l1, l2

# Model architecture visualization
from visualkeras import layered_view

# Model training
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.metrics import Precision, Recall, AUC

# Make plots have guidelines
plt.style.use('ggplot')

# Squash Python warnings
import warnings
warnings.filterwarnings('ignore')

```

```

WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1741311938.685213 1022400 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1741311938.690879 1022400 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has

```

already been registered

1.1 | Set Random Seed for Reproducibility

```
[2]: tf.keras.utils.set_random_seed(42)
      tf.random.set_seed(42)
      np.random.seed(42)
      random.seed(42)
```

1.2 | Declare Global Variables

```
[3]: # Define training and testing image directories
      home_dir = '/home/wiley'
      trn_dir = home_dir+'/regis/dataScience/kaggleProject/images/data/training'
      tst_dir = home_dir+'/regis/dataScience/kaggleProject/images/data/testing'

      # Define classes
      classes = ['negative', 'positive']

      # Image size and shape
      img_size = (224, 224)
      img_shape = (224, 224, 3)

      # Number of classes
      num_classes = 2

      # Declare batch size
      batch_size = 64

      # Flag to save weights
      save = True
```

2.0 | Define Functions

2.1 | Load DataFrames - Join image filename and path information - Create labels from class directory names - Create dataframe - Randomize dataframe rows

```
[4]: def load_dataframe(path):
      # Derive image file paths and labels from directory structure
      labels, paths = zip(*[(label, os.path.join(path, label, image))
                             for label in os.listdir(path)
                             if os.path.isdir(os.path.join(path, label))
                             for image in os.listdir(os.path.join(path, label))])

      # Create DataFrame
      df = pd.DataFrame({'paths': paths, 'labels': labels})

      # Randomize rows to help eliminate bias
```

```
df = df.sample(frac=1, random_state=42).reset_index(drop=True)

return df
```

2.2 | Plot Performance Metrics Plot the following: - Training loss - Validation loss - Training Accuracy - Validation Accuracy - Training Precision - Validation Precision - Training Recall - Validation Recall - Training AUC - Validation AUC

```
[5]: def plot_history(history):
    epochs = range(1, len(history.history['accuracy']) + 1)

    # Plot training and validation loss
    plt.figure(figsize=(20,12))
    plt.subplot(2,2,1)
    plt.plot(epochs, history.history['loss'], 'b', label = 'Training Loss')
    plt.plot(epochs, history.history['val_loss'], 'r', label = 'Validation_
↳Loss')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()

    # Plot training and validation accuracy
    plt.subplot(2,2,2)
    plt.plot(epochs, history.history['accuracy'], 'b', label = 'Training_
↳Accuracy')
    plt.plot(epochs, history.history['val_accuracy'], 'r', label = 'Validation_
↳Accuracy')
    plt.title('Training and Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()

    plt.suptitle('Model Loss and Accuracy over Epochs', fontsize=16)
    plt.show()

    # Plot training and validation precision
    plt.figure(figsize=(20,12))
    plt.subplot(2,2,1)
    plt.plot(epochs, history.history['precision'], 'b', label='Training_
↳Precision')
    plt.plot(epochs, history.history['val_precision'], 'r', label='Validation_
↳Precision')
    plt.title('Training and Validation Precision')
    plt.xlabel('Epochs')
    plt.ylabel('Precision')
    plt.legend()
```

```

# Plot training and validation recall
plt.subplot(2,2,2)
plt.plot(epochs, history.history['recall'], 'b', label='Training Recall')
plt.plot(epochs, history.history['val_recall'], 'r', label='Validation_
↪Recall')
plt.title('Training and Validation Recall')
plt.xlabel('Epochs')
plt.ylabel('Recall')
plt.legend()

plt.suptitle('Model Precision and Recall over Epochs', fontsize=16)
plt.show()

# Plot training and validation AUC
plt.figure(figsize=(8,6))
plt.plot(epochs, history.history['auc'], 'b', label='Training AUC')
plt.plot(epochs, history.history['val_auc'], 'r', label='Validation AUC')
plt.title('Training and Validation AUC')
plt.xlabel('Epochs')
plt.ylabel('Recall')
plt.legend()
plt.show()

```

2.3 | Evaluate Model's Performance on Test DataSet - Infer loss, accuracy, precision, recall, and AUC from dataset - Compute F1 Score from precision and recall

```

[6]: def score_model(model, ds):
    # Get metrics from test data
    loss, acc, auc, prec, recall = model.evaluate(ds)

    # Calculate F1 Score from precision and recall
    f1_score = 2 * (prec * recall) / (prec + recall)

    # Print results
    print('-' * 30)
    print(f'Loss:      {loss:.4f}')
    print(f'Accuracy:   {acc:.4f}')
    print(f'Precision:   {prec:.4f}')
    print(f'Recall:     {recall:.4f}')
    print(f'AUC:       {auc:.4f}')
    print(f'F1 Score:    {f1_score:.4f}')
    print('-' * 30)

```

2.4 | Plot Confusion Matrix

```

[7]: def plot_cm(model, ds):
    # Get predictions from dataset

```

```

preds = np.argmax(np.round(model.predict(ds)), axis=1)

# Create confusion matrix
cm = confusion_matrix(ds.classes, preds)

# Visualize confusion matrix
plt.figure(figsize=(5,3))
sns.heatmap(cm, annot=True, fmt='d', cmap='viridis',
            xticklabels=classes,
            yticklabels=classes)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```

2.5 | Compute TPR and TNR

```

[8]: def compute_tpr(model, ds):
    # get predictions from dataset
    preds = np.argmax(np.round(model.predict(ds)), axis=1)

    # Create confusion matrix
    cm = confusion_matrix(ds.classes, preds)

    # Extract required values from confusion matrix
    (tn, fp, fn, tp) = cm.flatten()

    # Calculate TPR
    tpr = tp / (tp + fn)

    # Calculate TNR
    tnr = tn / (tn + fp)

    # Print TPR and TNR
    print('-' * 30)
    print(f'True Positive Rate (TPR): {tpr:.4f}')
    print(f'True Negative Rate (TNR): {tnr:.4f}')
    print('-' * 30)

```

2.6 | Show True vs Predicted Labels

```

[9]: def display_preds(model, ds):
    # Extract true and predicted labels from dataset
    images, labels = next(ds)
    preds = model.predict(images)
    pred_labs = np.argmax(preds, axis=1)
    dict = ds.class_indices
    tr_labels = list(dict.keys())

```

```

# Plot the images with true and predicted labels
plt.figure(figsize=(20,20))
for i in range(16):
    img = images[i]
    label = labels[i]
    tr_label = classes[np.argmax(label)]
    pred_label = classes[pred_labs[i]]
    plt.subplot(4,4,i+1)
    plt.imshow(img)
    plt.title('True:      %s \nPredict:  %s' % (tr_label, pred_label),
    ↪fontsize=15,
    loc='left', color='k')
    plt.axis('off')
plt.tight_layout()
plt.show()

```

2.7 | Print Images

```

[10]: def print_images(ds):
# Pull images and labels out of the dataset
images, labels = next(ds)

# Create a dictionary of class indices
dict = ds.class_indices

# Form classes from the dictionary created in last step
classes = list(dict.keys())

# Plot the images and labels -- 16 images at a time
plt.figure(figsize=(20,20))
for i in range(16):
    img = images[i]
    label = labels[i]
    class_name = classes[np.argmax(label)]
    plt.subplot(4,4,i+1)
    plt.imshow(img)
    plt.title(class_name, loc='left', fontsize=15)
    plt.axis('off')

plt.show()

```

3.0 | Load Data

3.1 | Create and Load DataFrame for EDA

```
[11]: # Load training data
trn_df = load_dataframe(trn_dir)

# Load testing data
tst_df = load_dataframe(tst_dir)

# Take a look at the results
print('Training:  \n', trn_df.head(10).to_markdown())
print('Testing:   \n', tst_df.head(10).to_markdown())
```

```
Training:
|   | paths
| labels |
|---:|:-----|
| 0 | /home/wiley/regis/dataScience/kaggleProject/images/data/training/negative
/image(192).jpg | negative |
| 1 | /home/wiley/regis/dataScience/kaggleProject/images/data/training/negative
/image(122).jpg | negative |
| 2 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (769).jpg | negative |
| 3 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (511).jpg | negative |
| 4 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer
(720).jpg | positive |
| 5 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (468).jpg | negative |
| 6 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (1108).jpg | negative |
| 7 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (360).jpg | negative |
| 8 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer
(402).jpg | positive |
| 9 |
/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not
Cancer (177).jpg | negative |
Testing:
|   | paths
| labels |
|---:|:-----|
| 0 | /home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative
/image(192).jpg | negative |
| 1 | /home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative
/image(122).jpg | negative |
| 2 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (769).jpg | negative |
| 3 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (511).jpg | negative |
| 4 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(720).jpg | positive |
| 5 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (468).jpg | negative |
| 6 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (1108).jpg | negative |
| 7 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (360).jpg | negative |
| 8 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(402).jpg | positive |
| 9 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (177).jpg | negative |
```



```

| 0 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(1404).jpg      | positive |
| 1 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (1168).jpg | negative |
| 2 | /home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/
image(30).jpg      | negative |
| 3 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (720).jpg  | negative |
| 4 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(1559).jpg      | positive |
| 5 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(2272).jpg      | positive |
| 6 | /home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/
image(33).jpg      | negative |
| 7 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(714).jpg      | positive |
| 8 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/negative/Not
Cancer (1090).jpg | negative |
| 9 |
/home/wiley/regis/dataScience/kaggleProject/images/data/testing/positive/Cancer
(300).jpg      | positive |

## 4.0 | EDA


```

4.1 | Look at Training Images' Distribution

```

[12]: plt.figure(figsize=(6,4))
      trn_df['labels'].value_counts().plot(kind='bar')
      plt.title('Distribution of Image Counts in Training Data')
      plt.xlabel('Category')
      plt.ylabel('Image Count')
      plt.show()


```



Negative images slightly outnumber the positive ones, but are close enough to continue without additional data wrangling

4.2 | Look at Testing Images' Distribution

```
[13]: plt.figure(figsize=(6,4))
      tst_df['labels'].value_counts().plot(kind='bar')
      plt.title('Distribution of Image Counts in Testing Data')
      plt.xlabel('Category')
      plt.ylabel('Image Count')
      plt.show()
```



Distribution mirrors what the *training data* shows, but with less frequency.

4.3 | Examine Shape of Training and Testing DataFrames

```
[14]: print('Training Shape: \n', trn_df.shape)
      print('Testing Shape: \n', tst_df.shape)
```

Training Shape:

(4076, 2)

Testing Shape:

(1024, 2)

NOTE: Since the dataframes are built from the contents of the image directories, there should be no missing values or duplicates.

4.0 | Data Wrangling

4.1 | Create a Validation Subset from Training Data I will use `flow_from_dataframe()` to create datasets for model training; therefore, no reason to create a new directory structure for validation data

```
[15]: val_df, trn_df = train_test_split(trn_df, train_size=0.2, random_state=42,
                                     stratify=trn_df['labels'])
print(val_df.sample(10).to_markdown())
print(f'Validation Shape: {val_df.shape}')
```

	paths	labels
2341	/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer (404).jpg	positive
1615	/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer (830).jpg	positive
3949	/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not Cancer (595).jpg	negative
2158	/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer (1965).jpg	positive
51	/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer (1479).jpg	positive
1967	/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not Cancer (1274).jpg	negative
3644	/home/wiley/regis/dataScience/kaggleProject/images/data/training/positive/Cancer (2414).jpg	positive
2917	/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not Cancer (546).jpg	negative
1429	/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not Cancer (225).jpg	negative
2855	/home/wiley/regis/dataScience/kaggleProject/images/data/training/negative/Not Cancer (1673).jpg	negative

Validation Shape: (815, 2)

4.2 | Process Images from DataFrames Image augmentation will be used on the training and validation datasets. The test images will just be normalized.

```
[16]: # Apply image augmentation
gen = ImageDataGenerator(rescale=1./255,
                        brightness_range=(0.5, 1.5),
                        rotation_range=20,
```

```

        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2)

# The test dataset should not be augmented
# just rescaled
tst_gen = ImageDataGenerator(rescale=1./255)

# Create training datagen set
trn_gen = gen.flow_from_dataframe(trn_df, x_col='paths', y_col='labels',
                                batch_size=batch_size, target_size=img_size,
                                shuffle=True)

# Create validation datagen set
val_gen = gen.flow_from_dataframe(val_df, x_col='paths', y_col='labels',
                                batch_size=batch_size, target_size=img_size,
                                shuffle=True)

# Create test datagen set
tst_gen = tst_gen.flow_from_dataframe(tst_df, x_col='paths', y_col='labels',
                                    batch_size=16, target_size=img_size,
                                    shuffle=False)

```

Found 3261 validated image filenames belonging to 2 classes.

Found 815 validated image filenames belonging to 2 classes.


Found 1024 validated image filenames belonging to 2 classes.

4.3 | Examine a few Augmented Images and their Labels The images displayed have been augmented in the previous step. The appearance may not be consistent with non-augmented images.

```

[17]: # Print augmented images from training dataset
      print_images(trn_gen)

```



5.0 | Configure Training Values

5.1 | Basic Values

```
[18]: # Number of training epochs
epochs = 50

# Steps per epoch
steps_per_ep = trn_gen.samples // batch_size

# Validation steps
val_steps = val_gen.samples // batch_size
```

```

print(f'Image shape:      {img_shape}')
print(f'Epochs:          {epochs}')
print(f'Batch size:        {batch_size}')
print(f'Steps per epoch:    {steps_per_ep}')
print(f'Validation steps:   {val_steps}')

```

```

Image shape:      (224, 224, 3)
Epochs:          50
Batch size:        64
Steps per epoch:  50
Validation steps:  12

```

5.2 | Define Callbacks With these *callbacks* the model's training will stop if the training loss stops decreasing (EarlyStopping()), and the learning rate will be reduced until the validation loss plateaus (ReduceLROnPlateau())

```

[19]: # Define early_stop callback
early_stop = EarlyStopping(monitor='loss', min_delta=0.000000001, patience=5,
                           baseline=None, restore_best_weights=True,
                           ↪start_from_epoch=0)

# Define reduce LR on Plateau callback
reduceLR0 = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=8,
                               ↪mode='auto',
                               min_delta=0.0001, cooldown=0, min_lr=0)

```

6.0 | Baseline Model ### Define Model's Architecture

6.1 | Define Model's Architecture The CNN model is being defined by using `models.Sequential()` method. It consists of four convolution layers flattened into two fully connected layers with dropout. The output layer will use the *softmax* activation function instead of *relu*

```

[20]: backend.clear_session()

model_cnn = models.Sequential([
    # Conv layer #1
    Conv2D(32, (4,4), activation='relu', input_shape=img_shape),
    MaxPooling2D(pool_size=(3,3)),

    # Conv layer #2
    Conv2D(64, (4,4), activation='relu'),
    MaxPooling2D(pool_size=(3,3)),

    # Conv layer #3
    Conv2D(128, (4,4), activation='relu'),

```

```

MaxPooling2D(pool_size=(4,4)),

# Conv layer #4
Conv2D(128, (4,4), activation='relu'),
Flatten(),

# Fully connect layers
Dense(512, activation='relu'),
Dropout(0.5, seed=42),
Dense(num_classes, activation='softmax')
])

model_cnn.summary()

```

I0000 00:00:1741311942.601514 1022400 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 9564 MB memory: -> device: 0, name: NVIDIA GeForce RTX 4070, pci bus id: 0000:09:00.0, compute capability: 8.9

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 221, 221, 32)	1,568
max_pooling2d (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_1 (Conv2D)	(None, 70, 70, 64)	32,832
max_pooling2d_1 (MaxPooling2D)	(None, 23, 23, 64)	0
conv2d_2 (Conv2D)	(None, 20, 20, 128)	131,200
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 128)	0
conv2d_3 (Conv2D)	(None, 2, 2, 128)	262,272
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 512)	262,656
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 2)	1,026

Total params: 691,554 (2.64 MB)


Trainable params: 691,554 (2.64 MB)

Non-trainable params: 0 (0.00 B)

6.2 | Visualize Layers

```
[21]: layered_view(model_cnn, legend=True, max_xy=300)
```

[21]:



6.3 | Compile and Train Model The `Adam()` optimizer was selected for this model, since it is well suited to classification problems. The loss function `categorical_crossentropy()` was also selected for the same reason.

```
[22]: # Configure Adam optimizer
opt = optimizers.Adam(learning_rate=0.001, beta_1=0.869, beta_2=0.995)
```

```

# Compile base model
model_cnn.compile(optimizer=opt, loss='categorical_crossentropy',
                  metrics=['accuracy', tf.keras.metrics.
    ↳Precision(name='precision'),
                                tf.keras.metrics.Recall(name='recall'),
                                tf.keras.metrics.AUC(curve='PR', name='auc')]))

# Fit data to model and record training history
hist_cnn = model_cnn.fit(trn_gen, batch_size=batch_size,
    ↳steps_per_epoch=steps_per_ep,
                        epochs=epochs, validation_data=val_gen,
                        validation_steps=val_steps,
                        callbacks=[early_stop, reduceLRO])

```

Epoch 1/50

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1741311946.973225 1022525 service.cc:148] XLA service 0x7f482c00e2a0 initialized for platform CUDA (this does not guarantee that XLA will be used).

Devices:

I0000 00:00:1741311946.973248 1022525 service.cc:156] StreamExecutor device (0): NVIDIA GeForce RTX 4070, Compute Capability 8.9

I0000 00:00:1741311947.266275 1022525 cuda_dnn.cc:529] Loaded cuDNN version 90300

2/50 1s 29ms/step - accuracy: 0.5586 - auc: 0.5534 - loss: 0.6881 - precision: 0.5586 - recall: 0.5586

I0000 00:00:1741311952.333143 1022525 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

50/50 53s 918ms/step - accuracy: 0.5327 - auc: 0.5350 - loss: 0.6916 - precision: 0.5327 - recall: 0.5327 - val_accuracy: 0.6445 - val_auc: 0.7139 - val_loss: 0.6248 - val_precision: 0.6445 - val_recall: 0.6445 - learning_rate: 0.0010

Epoch 2/50

50/50 2s 47ms/step - accuracy: 0.5781 - auc: 0.6129 - loss: 0.6792 - precision: 0.5781 - recall: 0.5781 - val_accuracy: 0.6809 - val_auc: 0.6446 - val_loss: 0.6550 - val_precision: 0.6809 - val_recall: 0.6809 - learning_rate: 0.0010

Epoch 3/50

50/50 36s 719ms/step - accuracy: 0.6224 - auc: 0.6676 - loss: 0.6430 - precision: 0.6224 - recall: 0.6224 - val_accuracy: 0.7214 - val_auc: 0.7645 - val_loss: 0.5792 - val_precision: 0.7214 - val_recall: 0.7214 - learning_rate: 0.0010

Epoch 4/50

50/50 0s 8ms/step -

accuracy: 0.7812 - auc: 0.8293 - loss: 0.5304 - precision: 0.7812 - recall:
 0.7812 - val_accuracy: 0.7234 - val_auc: 0.8020 - val_loss: 0.5543 -
 val_precision: 0.7234 - val_recall: 0.7234 - learning_rate: 0.0010
 Epoch 5/50
 50/50 37s 752ms/step -
 accuracy: 0.6993 - auc: 0.7528 - loss: 0.5854 - precision: 0.6993 - recall:
 0.6993 - val_accuracy: 0.7018 - val_auc: 0.7576 - val_loss: 0.5844 -
 val_precision: 0.7018 - val_recall: 0.7018 - learning_rate: 0.0010
 Epoch 6/50
 50/50 0s 8ms/step -
 accuracy: 0.7031 - auc: 0.7551 - loss: 0.5706 - precision: 0.7031 - recall:
 0.7031 - val_accuracy: 0.7021 - val_auc: 0.7787 - val_loss: 0.5795 -
 val_precision: 0.7021 - val_recall: 0.7021 - learning_rate: 0.0010
 Epoch 7/50
 50/50 37s 739ms/step -
 accuracy: 0.7264 - auc: 0.7952 - loss: 0.5436 - precision: 0.7264 - recall:
 0.7264 - val_accuracy: 0.7474 - val_auc: 0.8068 - val_loss: 0.5323 -
 val_precision: 0.7474 - val_recall: 0.7474 - learning_rate: 0.0010
 Epoch 8/50
 50/50 0s 9ms/step -
 accuracy: 0.8281 - auc: 0.9322 - loss: 0.4185 - precision: 0.8281 - recall:
 0.8281 - val_accuracy: 0.7021 - val_auc: 0.8342 - val_loss: 0.5065 -
 val_precision: 0.7021 - val_recall: 0.7021 - learning_rate: 0.0010
 Epoch 9/50
 50/50 37s 737ms/step -
 accuracy: 0.7549 - auc: 0.8348 - loss: 0.4961 - precision: 0.7549 - recall:
 0.7549 - val_accuracy: 0.7370 - val_auc: 0.8335 - val_loss: 0.5057 -
 val_precision: 0.7370 - val_recall: 0.7370 - learning_rate: 0.0010
 Epoch 10/50
 50/50 0s 9ms/step -
 accuracy: 0.8281 - auc: 0.8995 - loss: 0.4141 - precision: 0.8281 - recall:
 0.8281 - val_accuracy: 0.5745 - val_auc: 0.7324 - val_loss: 0.6385 -
 val_precision: 0.5745 - val_recall: 0.5745 - learning_rate: 0.0010
 Epoch 11/50
 50/50 37s 737ms/step -
 accuracy: 0.7708 - auc: 0.8366 - loss: 0.4923 - precision: 0.7708 - recall:
 0.7708 - val_accuracy: 0.7904 - val_auc: 0.8898 - val_loss: 0.4197 -
 val_precision: 0.7904 - val_recall: 0.7904 - learning_rate: 0.0010
 Epoch 12/50
 50/50 0s 8ms/step -
 accuracy: 0.7500 - auc: 0.8306 - loss: 0.5260 - precision: 0.7500 - recall:
 0.7500 - val_accuracy: 0.8511 - val_auc: 0.9182 - val_loss: 0.3815 -
 val_precision: 0.8511 - val_recall: 0.8511 - learning_rate: 0.0010
 Epoch 13/50
 50/50 37s 747ms/step -
 accuracy: 0.8001 - auc: 0.8799 - loss: 0.4332 - precision: 0.8001 - recall:
 0.8001 - val_accuracy: 0.7396 - val_auc: 0.8235 - val_loss: 0.5574 -
 val_precision: 0.7396 - val_recall: 0.7396 - learning_rate: 0.0010

Epoch 14/50
50/50 0s 8ms/step -
accuracy: 0.7812 - auc: 0.8671 - loss: 0.4765 - precision: 0.7812 - recall:
0.7812 - val_accuracy: 0.7021 - val_auc: 0.7781 - val_loss: 0.6109 -
val_precision: 0.7021 - val_recall: 0.7021 - learning_rate: 0.0010

Epoch 15/50
50/50 37s 741ms/step -
accuracy: 0.8138 - auc: 0.9012 - loss: 0.3980 - precision: 0.8138 - recall:
0.8138 - val_accuracy: 0.7904 - val_auc: 0.8892 - val_loss: 0.4228 -
val_precision: 0.7904 - val_recall: 0.7904 - learning_rate: 0.0010

Epoch 16/50
50/50 0s 8ms/step -
accuracy: 0.7812 - auc: 0.8906 - loss: 0.4217 - precision: 0.7812 - recall:
0.7812 - val_accuracy: 0.7660 - val_auc: 0.8260 - val_loss: 0.5120 -
val_precision: 0.7660 - val_recall: 0.7660 - learning_rate: 0.0010

Epoch 17/50
50/50 37s 734ms/step -
accuracy: 0.8042 - auc: 0.8943 - loss: 0.4100 - precision: 0.8042 - recall:
0.8042 - val_accuracy: 0.8346 - val_auc: 0.9057 - val_loss: 0.3876 -
val_precision: 0.8346 - val_recall: 0.8346 - learning_rate: 0.0010

Epoch 18/50
50/50 0s 8ms/step -
accuracy: 0.8594 - auc: 0.9460 - loss: 0.3123 - precision: 0.8594 - recall:
0.8594 - val_accuracy: 0.8085 - val_auc: 0.8891 - val_loss: 0.4221 -
val_precision: 0.8085 - val_recall: 0.8085 - learning_rate: 0.0010

Epoch 19/50
50/50 37s 741ms/step -
accuracy: 0.8420 - auc: 0.9185 - loss: 0.3623 - precision: 0.8420 - recall:
0.8420 - val_accuracy: 0.8451 - val_auc: 0.9014 - val_loss: 0.3982 -
val_precision: 0.8451 - val_recall: 0.8451 - learning_rate: 0.0010

Epoch 20/50
50/50 0s 8ms/step -
accuracy: 0.8594 - auc: 0.9371 - loss: 0.3294 - precision: 0.8594 - recall:
0.8594 - val_accuracy: 0.8085 - val_auc: 0.8938 - val_loss: 0.4193 -
val_precision: 0.8085 - val_recall: 0.8085 - learning_rate: 0.0010

Epoch 21/50
50/50 38s 760ms/step -
accuracy: 0.8514 - auc: 0.9304 - loss: 0.3381 - precision: 0.8514 - recall:
0.8514 - val_accuracy: 0.8333 - val_auc: 0.9204 - val_loss: 0.3603 -
val_precision: 0.8333 - val_recall: 0.8333 - learning_rate: 1.0000e-04

Epoch 22/50
50/50 0s 9ms/step -
accuracy: 0.8594 - auc: 0.9448 - loss: 0.3159 - precision: 0.8594 - recall:
0.8594 - val_accuracy: 0.8936 - val_auc: 0.9403 - val_loss: 0.3106 -
val_precision: 0.8936 - val_recall: 0.8936 - learning_rate: 1.0000e-04

Epoch 23/50
50/50 37s 744ms/step -
accuracy: 0.8650 - auc: 0.9418 - loss: 0.3083 - precision: 0.8650 - recall:

```

0.8650 - val_accuracy: 0.8711 - val_auc: 0.9370 - val_loss: 0.3210 -
val_precision: 0.8711 - val_recall: 0.8711 - learning_rate: 1.0000e-04
Epoch 24/50
50/50          0s 8ms/step -
accuracy: 0.8594 - auc: 0.9278 - loss: 0.3420 - precision: 0.8594 - recall:
0.8594 - val_accuracy: 0.8511 - val_auc: 0.9113 - val_loss: 0.3816 -
val_precision: 0.8511 - val_recall: 0.8511 - learning_rate: 1.0000e-04
Epoch 25/50
50/50          37s 745ms/step -
accuracy: 0.8750 - auc: 0.9467 - loss: 0.2970 - precision: 0.8750 - recall:
0.8750 - val_accuracy: 0.8685 - val_auc: 0.9407 - val_loss: 0.3157 -
val_precision: 0.8685 - val_recall: 0.8685 - learning_rate: 1.0000e-04
Epoch 26/50
50/50          0s 8ms/step -
accuracy: 0.8594 - auc: 0.9741 - loss: 0.2463 - precision: 0.8594 - recall:
0.8594 - val_accuracy: 0.8511 - val_auc: 0.9108 - val_loss: 0.3802 -
val_precision: 0.8511 - val_recall: 0.8511 - learning_rate: 1.0000e-04
Epoch 27/50
50/50          37s 740ms/step -
accuracy: 0.8809 - auc: 0.9519 - loss: 0.2838 - precision: 0.8809 - recall:
0.8809 - val_accuracy: 0.8607 - val_auc: 0.9347 - val_loss: 0.3223 -
val_precision: 0.8607 - val_recall: 0.8607 - learning_rate: 1.0000e-04
Epoch 28/50
50/50          0s 9ms/step -
accuracy: 0.8750 - auc: 0.9548 - loss: 0.2916 - precision: 0.8750 - recall:
0.8750 - val_accuracy: 0.8936 - val_auc: 0.9668 - val_loss: 0.2490 -
val_precision: 0.8936 - val_recall: 0.8936 - learning_rate: 1.0000e-04
Epoch 29/50
50/50          37s 752ms/step -
accuracy: 0.8945 - auc: 0.9569 - loss: 0.2696 - precision: 0.8945 - recall:
0.8945 - val_accuracy: 0.8763 - val_auc: 0.9441 - val_loss: 0.3052 -
val_precision: 0.8763 - val_recall: 0.8763 - learning_rate: 1.0000e-04
Epoch 30/50
50/50          0s 8ms/step -
accuracy: 0.8750 - auc: 0.9667 - loss: 0.2531 - precision: 0.8750 - recall:
0.8750 - val_accuracy: 0.9149 - val_auc: 0.9275 - val_loss: 0.3247 -
val_precision: 0.9149 - val_recall: 0.9149 - learning_rate: 1.0000e-04
Epoch 31/50
50/50          36s 727ms/step -
accuracy: 0.8791 - auc: 0.9509 - loss: 0.2869 - precision: 0.8791 - recall:
0.8791 - val_accuracy: 0.8516 - val_auc: 0.9363 - val_loss: 0.3248 -
val_precision: 0.8516 - val_recall: 0.8516 - learning_rate: 1.0000e-04

## 7.0 | Evaluate Performance

```


7.1 | Plot Training and Validation Metrics If the training and validation metrics diverge significantly from each other, that can be an indication of overfitting.


```
[23]: plot_history(hist_cnn)
```

Model Loss and Accuracy over Epochs



Model Precision and Recall over Epochs





7.2 | Score Model To evaluate the model's performance the following matrices will be evaluated against the test dataset: - Model Loss — gives a nuanced view of model optimization - Model Accuracy — provides the proportion of all classifications that were correct - Recall — proportion of correct positive classifications - Area Under Curve (AUC) — represents the probability that the model, if given a randomly chosen positive and negative example, will rank the positive higher than the negative - F1 Score — describes the harmonic mean of the precision and recall of the model

```
[24]: score_model(model_cnn, tst_gen)
```

```
64/64          3s 25ms/step -
accuracy: 0.8804 - auc: 0.9455 - loss: 0.2897 - precision: 0.8804 - recall:
0.8804
```


```
-----
Loss:      0.2621
Accuracy:  0.8916
Precision: 0.8916
Recall:    0.9568
AUC:       0.8916
F1 Score:  0.9231
-----
```

7.3 | Plot Confusion Matrix A confusion matrix provides a visual representation of a model's performance when it comes to comparing true positives, false negatives, true negatives, and false positives.

```
[25]: plot_cm(model_cnn, tst_gen)
```

64/64

2s 24ms/step



7.4 | Compute TPR and TNR The True Positive Rate (TPR) and True Negative Rate (TNR) are good indicators of how well the model is predicting positives (1s) and negatives (0s).

```
[26]: compute_tpr(model_cnn, tst_gen)
```

64/64

2s 25ms/step

True Positive Rate (TPR): 0.9143


True Negative Rate (TNR): 0.8697

7.5 | View and Compare Predicted with True Labels If the *True* and *Predict* labels match, then the model accurately predicted the label of the image

```
[27]: display_preds(model_cnn, tst_gen)
```

1/1

1s 652ms/step



8.0 | Save Weights

```
[28]: # Use date and time to create a unique filename
if save:
    now = datetime.now()
    date = now.strftime('%Y%m%d')
    time = now.strftime('%H%M%S')
    filename = date+time+'_cnn_brain_tumor.keras'

    # Save weights in keras zip format
    model_cnn.save('../weights/'+filename)
```

9.0 | Discussion and Conclusions

9.1 | General Performance The model is generalizing the test data well. The plots indicated the model is training well and overfitting is minimal. Training and Validation metrics closely follow each other during the training process and I do not see any items that are concerning. F1, True Positive, and True Negative scores are above 0.90 which indicates it is accurately predicting positive and negative labels with 90% accuracy.

9.2 | Further Work While the model is generalizing the test data with acceptable results, it can do better. I will conduct further research on CNN model tuning and keep experimenting until results are consistently in the high 90s.

[]: